

In [2]:

```
%matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
```

In [3]:

```
data = pd.read_csv('health_care_diabetes_raw.csv')
```

In [4]:

```
data.head()
```

Out[4]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

In [5]:

```
data.shape
```

Out[5]:

(768, 9)

Project Task: Week 1 -- Data Exploration and Missing Values Treatment

In [6]:

```
#Checking for null values in Dataset
data.isnull().any()
```

Out[6]:

```
Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction  False
Age               False
Outcome           False
dtype: bool
```

Since the 0 value in Glucose,BloodPressure,SkinThickness,Insulin and BMI variables represent missing values.Lets find now many instances are there in each of the above variables

In [7]:

```
data[data['Glucose']==0]
```

Out[7]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
75	1	0	48	20	0	24.7	0.140	22	0
182	1	0	74	20	23	27.7	0.299	21	0
342	1	0	68	35	0	32.0	0.389	22	0
349	5	0	80	32	0	41.0	0.346	37	1
502	6	0	68	41	0	39.0	0.727	41	1

In [8]:

```
(5/765)*100
#only 0.6% of data is having missing values in Glucose column. No need to worry we can ignore them
```

Out[8]:

0.6535947712418301

In [9]:

```
(data[data['BloodPressure']==0]).shape
```

Out[9]:

(35, 9)

In [10]:

```
(35/765)*100  
#4.5% of data is having missing values in BloodPressure column
```

Out[10]:

4.57516339869281

In [11]:

```
(data[data['SkinThickness']==0]).shape
```

Out[11]:

(227, 9)

In [13]:

```
(227/765)*100  
#29.6% of data is having missing values in SkinThickness column
```

Out[13]:

29.673202614379086

In [29]:

```
(data[data['Insulin']==0]).shape
```

Out[29]:

(0, 9)

In [30]:

```
(374/765)*100  
#~49% of data is having missing values in Insulin column
```

Out[30]:

48.888888888888886

In [31]:

```
(data[data['BMI']==0]).shape
```

Out[31]:

(11, 9)

In [17]:

```
(11/765)*100  
#1.4% of data is having missing values in BMI column
```

Out[17]:

1.4379084967320261


Since Insulin and SkinThickness are having higher percentages of missing values lets try to fill up the missing values

In [32]:

```
plt.hist(data['SkinThickness'],edgecolor='red')
```

Out[32]:

```
(array([ 59., 141., 408., 118., 36., 4., 1., 0., 0., 1.]),
 array([ 7. , 16.2, 25.4, 34.6, 43.8, 53. , 62.2, 71.4, 80.6, 89.8, 99. ]),
 <BarContainer object of 10 artists>)
```



In [33]:

```
data[data['SkinThickness']!=0]['SkinThickness'].describe()
```

Out[33]:


```
count    768.000000
mean      29.153420
std        8.790942
min         7.000000
25%       25.000000
50%       29.153420
75%       32.000000
max       99.000000
Name: SkinThickness, dtype: float64
```

In [34]:

```
plt.hist(data['Insulin'],edgecolor='red')
```

Out[34]:

```
(array([142., 517., 55., 29., 7., 10., 4., 1., 2., 1.]),
 array([ 14. , 97.2, 180.4, 263.6, 346.8, 430. , 513.2, 596.4, 679.6,
        762.8, 846. ]),
 <BarContainer object of 10 artists>)
```



In [35]:

```
data[data['Insulin']!=0]['Insulin'].describe()
```

Out[35]:

```
count    768.000000
mean    155.548223
std      85.021108
min      14.000000
25%     121.500000
50%     155.548223
75%     155.548223
max     846.000000
Name: Insulin, dtype: float64
```

Mean value of Skinthickness is ~29 and the mean value of Insulin is ~155 let impute the missing values with means


```
In [36]:
from numpy import nan
dataset_imputed = data
dataset_imputed[['SkinThickness','Insulin']] = dataset_imputed[['SkinThickness','Insulin']].replace(0, nan)
```

```
In [37]:
dataset_imputed.fillna(dataset_imputed.mean(), inplace=True)
```

```
In [38]:
plt.hist(dataset_imputed['Insulin'],edgecolor='red')
```

Out[38]:

```
(array([142., 517., 55., 29., 7., 10., 4., 1., 2., 1.]),
 array([ 14. , 97.2, 180.4, 263.6, 346.8, 430. , 513.2, 596.4, 679.6,
        762.8, 846. ]),
 <BarContainer object of 10 artists>)
```



```
In [39]:
data.describe()
```

Out[39]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	29.153420	155.548223	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	8.790942	85.021108	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	7.000000	14.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	25.000000	121.500000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	29.153420	155.548223	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	155.548223	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
In [40]:
dataset_imputed.describe()
```

Out[40]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	29.153420	155.548223	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	8.790942	85.021108	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	7.000000	14.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	25.000000	121.500000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	29.153420	155.548223	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	155.548223	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

In [42]:

```
dataset_imputed.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Pregnancies            768 non-null   int64  
1   Glucose                768 non-null   int64  
2   BloodPressure          768 non-null   int64  
3   SkinThickness          768 non-null   float64
4   Insulin                768 non-null   float64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64  
8   Outcome                768 non-null   int64  
dtypes: float64(4), int64(5)
memory usage: 54.1 KB
```

In [43]:

```
Positive = dataset_imputed[dataset_imputed['Outcome']==1]
Positive.head(5)
```

Out[43]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35.00000	155.548223	33.6	0.627	50	1
2	8	183	64	29.15342	155.548223	23.3	0.672	32	1
4	0	137	40	35.00000	168.000000	43.1	2.288	33	1
6	3	78	50	32.00000	88.000000	31.0	0.248	26	1
8	2	197	70	45.00000	543.000000	30.5	0.158	53	1

In [44]:

```
Negative = dataset_imputed[dataset_imputed['Outcome']==0]
Negative.head(5)
```

Out[44]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	1	85	66	29.00000	155.548223	26.6	0.351	31	0
3	1	89	66	23.00000	94.000000	28.1	0.167	21	0
5	5	116	74	29.15342	155.548223	25.6	0.201	30	0
7	10	115	0	29.15342	155.548223	35.3	0.134	29	0
10	4	110	92	29.15342	155.548223	37.6	0.191	30	0

In [53]:

```
dataset_imputed['Glucose'].value_counts().head(5)
```

Out[53]:


```
99    17
100    17
111    14
129    14
125    14
Name: Glucose, dtype: int64
```

In [54]:

```
plt.hist(dataset_imputed['Glucose'],edgecolor='red')
```

Out[54]:

```
(array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
 array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
        179.1, 199. ]),
 <BarContainer object of 10 artists>)
```



In [55]:

```
dataset_imputed['BloodPressure'].value_counts().head(7)
```

Out[55]:


```
70    57
74    52
78    45
68    45
72    44
64    43
80    40
Name: BloodPressure, dtype: int64
```

In [56]:

```
plt.hist(dataset_imputed['BloodPressure'],edgecolor='red')
```

Out[56]:

```
(array([ 35.,  1.,  2., 13., 107., 261., 243., 87., 14.,  5.]),
 array([ 0., 12.2, 24.4, 36.6, 48.8, 61., 73.2, 85.4, 97.6,
        109.8, 122. ]),
 <BarContainer object of 10 artists>)
```



In [57]:

```
dataset_imputed['SkinThickness'].value_counts().head(7)
```

Out[57]:


```
29.15342    227
32.00000     31
30.00000     27
27.00000     23
23.00000     22
33.00000     20
28.00000     20
Name: SkinThickness, dtype: int64
```

In [58]:

```
plt.hist(dataset_imputed['SkinThickness'],edgecolor='red')
```

Out[58]:

```
(array([ 59., 141., 408., 118., 36., 4., 1., 0., 0., 1.]),
 array([ 7., 16.2, 25.4, 34.6, 43.8, 53., 62.2, 71.4, 80.6, 89.8, 99. ]),
 <BarContainer object of 10 artists>)
```



In [59]:

```
dataset_imputed['Insulin'].value_counts().head(7)
```

Out[59]:


```
155.548223    374
105.000000     11
130.000000      9
140.000000      9
120.000000      8
94.000000       7
180.000000       7
Name: Insulin, dtype: int64
```

In [60]:

```
plt.hist(dataset_imputed['Insulin'],edgecolor='red')
```

Out[60]:

```
(array([142., 517., 55., 29., 7., 10., 4., 1., 2., 1.]),
 array([ 14., 97.2, 180.4, 263.6, 346.8, 430., 513.2, 596.4, 679.6,
        762.8, 846. ]),
 <BarContainer object of 10 artists>)
```



In [61]:

```
dataset_imputed['BMI'].value_counts().head(7)
```

Out[61]:


```
32.0    13
31.6    12
31.2    12
0.0     11
32.4    10
33.3    10
30.1     9
Name: BMI, dtype: int64
```

In [62]:

```
plt.hist(dataset_imputed['BMI'],edgecolor='red')
```

Out[62]:

(array([11., 0., 15., 156., 268., 224., 78., 12., 3., 1.]),
array([0. , 6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
 60.39, 67.1]),
<BarContainer object of 10 artists>)



In [63]:

```
dataset_imputed.describe().transpose()
```

Out[63]:

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.000000	6.000000	17.00
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.000000	140.250000	199.00
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.000000	80.000000	122.00
SkinThickness	768.0	29.153420	8.790942	7.000	25.00000	29.153420	32.000000	99.00
Insulin	768.0	155.548223	85.021108	14.000	121.50000	155.548223	155.548223	846.00
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.000000	36.600000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.372500	0.626250	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.000000	41.000000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.000000	1.000000	1.00

Project Task: Week 2 -- Corelation Analysis and Scatter Plots

In [64]:

```
Positive.shape
```

Out[64]:

(268, 9)

In [65]:

```
Negative.shape
```

Out[65]:


(500, 9)

In [66]:

```
plt.hist(Positive['BMI'],histtype='stepfilled',bins=20,edgecolor='red')
```

Out[66]:

```
(array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
        14.,  7.,  3.,  1.,  1.,  0.,  1.]),
 array([ 0.   ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
        26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
        53.68 , 57.035, 60.39 , 63.745, 67.1   ]),
 [<matplotlib.patches.Polygon at 0x258c736c940>])
```



In [67]:

```
Positive['BMI'].value_counts().head(7)
```

Out[67]:


```
32.9    8
31.6    7
33.3    6
31.2    5
30.5    5
32.0    5
34.3    4
Name: BMI, dtype: int64
```

In [68]:

```
plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20,edgecolor='red')
```

Out[68]:

```
(array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
        27., 29., 22., 24., 21., 25., 14.]),
 array([ 0.   ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
        159.2 , 169.15, 179.1 , 189.05, 199.   ]),
 [<matplotlib.patches.Polygon at 0x258c73df340>])
```



In [69]:

```
Positive['Glucose'].value_counts().head(7)
```

Out[69]:


```
125    7
128    6
129    6
115    6
158    6
146    5
124    5
Name: Glucose, dtype: int64
```

In [70]:

```
plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20,edgecolor='red')
```

Out[70]:

```
(array([16., 0., 0., 0., 0., 1., 0., 1., 6., 6., 19., 37., 56.,
        36., 41., 31., 7., 4., 4., 3.]),
 array([ 0.,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
        51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
        102.6, 108.3, 114. ]),
 [<matplotlib.patches.Polygon at 0x258c7440d60>])
```



In [72]:

```
Positive['BloodPressure'].value_counts().head(7)
```

Out[72]:


```
70    23
76    18
78    17
74    17
72    16
0      16
80    13
Name: BloodPressure, dtype: int64
```

In [73]:

```
plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20,edgecolor='red')
```

Out[73]:

```
(array([ 1.,  5., 11., 21., 113., 41., 34., 20., 15., 4., 1.,
        0.,  1., 0., 0., 0., 0., 0., 1.]),
 array([ 7. , 11.6, 16.2, 20.8, 25.4, 30. , 34.6, 39.2, 43.8, 48.4, 53. ,
        57.6, 62.2, 66.8, 71.4, 76. , 80.6, 85.2, 89.8, 94.4, 99. ]),
 [<matplotlib.patches.Polygon at 0x258c74a7d90>])
```



In [74]:

```
Positive['SkinThickness'].value_counts().head(7)
```

Out[74]:


```
29.15342    88
32.00000    14
30.00000     9
33.00000     9
39.00000     8
37.00000     8
36.00000     8
Name: SkinThickness, dtype: int64
```

In [75]:

```
plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20,edgecolor='red')
```

Out[75]:

```
(array([ 4., 12., 27., 169., 18., 10., 8., 5., 2., 1., 1.,
        6., 2., 1., 1., 0., 0., 0., 0., 1.]),
 array([ 14., 55.6, 97.2, 138.8, 180.4, 222., 263.6, 305.2, 346.8,
        388.4, 430., 471.6, 513.2, 554.8, 596.4, 638., 679.6, 721.2,
        762.8, 804.4, 846. ]),
 [<matplotlib.patches.Polygon at 0x258c750f3d0>])
```



In [76]:

```
Positive['Insulin'].value_counts().head(7)
```

Out[76]:

```
155.548223    138
130.000000     6
180.000000     4
175.000000     3
156.000000     3
185.000000     2
194.000000     2
Name: Insulin, dtype: int64
```

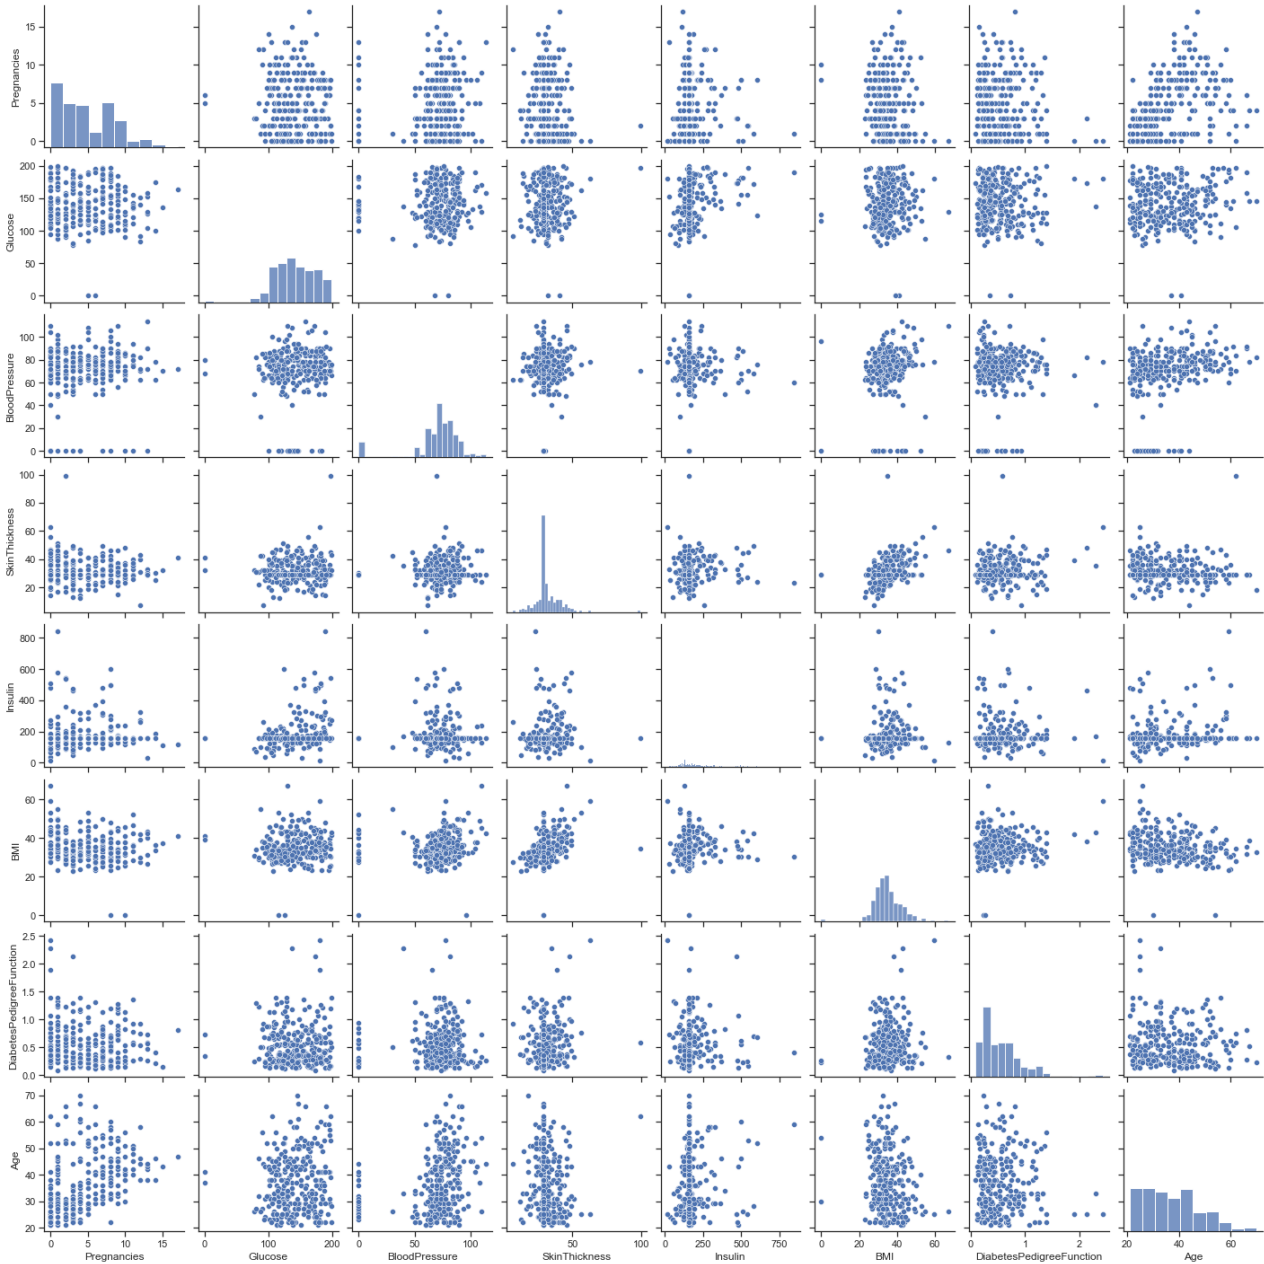
Scatter Plots

```
In [78]:  
  
#Pair plots for all dataset  
sns.set(style="ticks", color_codes=True)  
g = sns.pairplot(dataset_imputed,hue="Outcome")
```



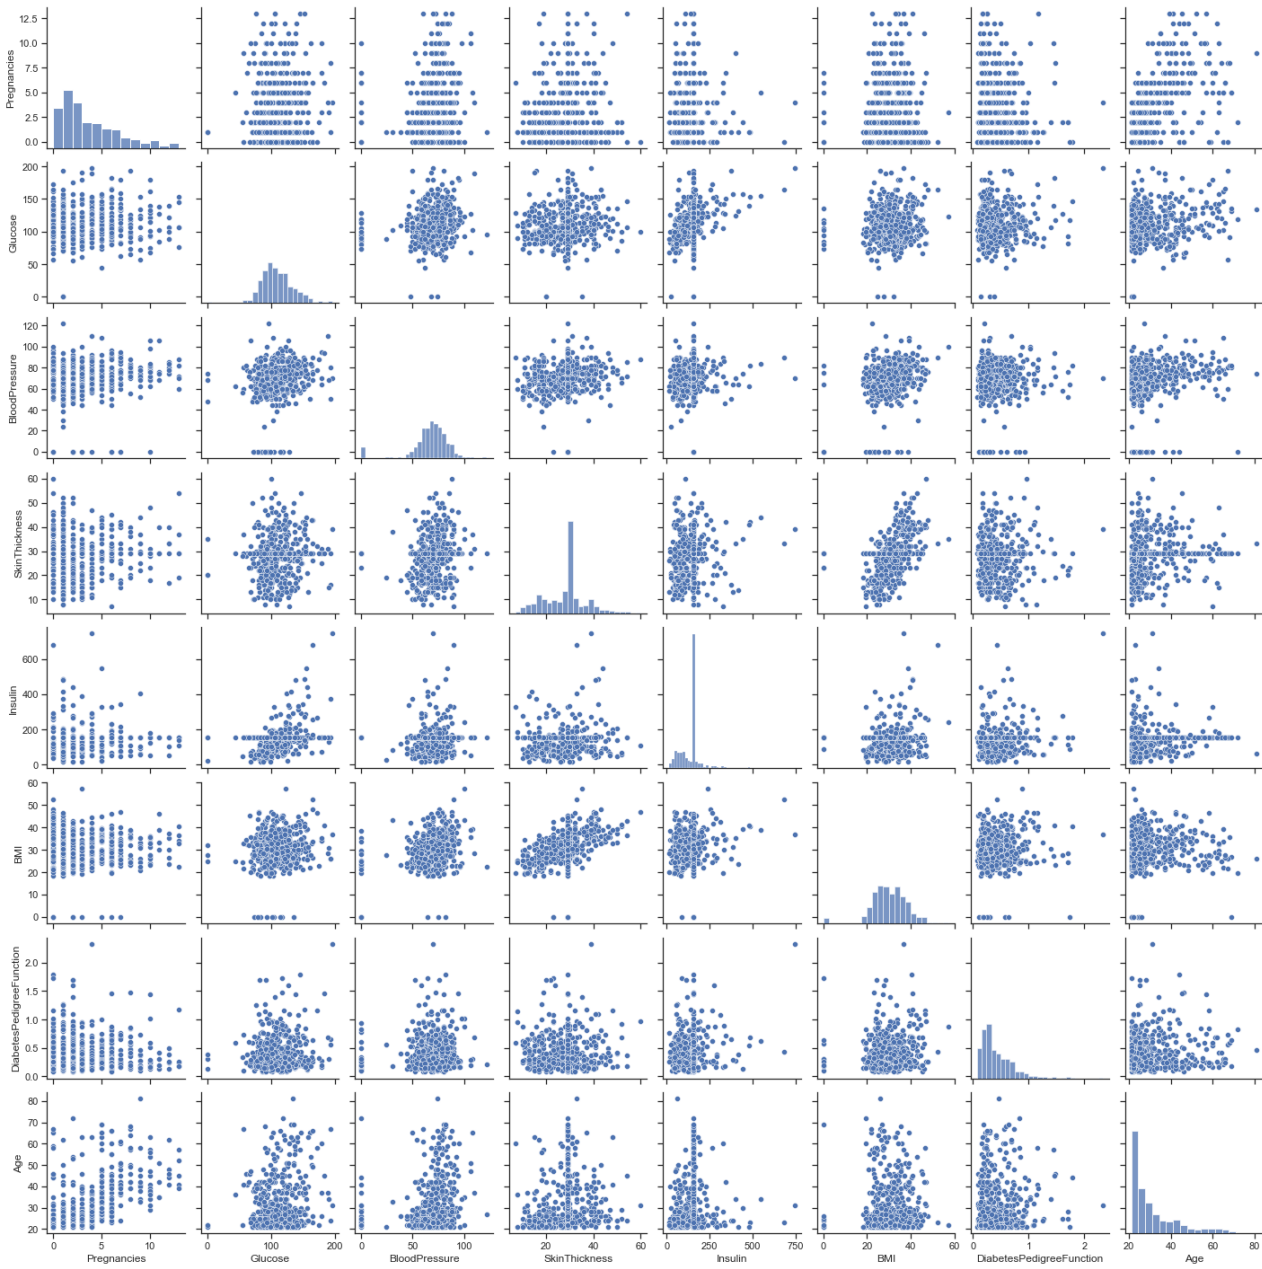
In [84]:

```
#Pair plots for all Positive cases
sns.set(style="ticks", color_codes=True)
g = sns.pairplot(Positive[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']])
```



In [86]:

```
#Pair plots for all Negative cases
sns.set(style="ticks", color_codes=True)
g = sns.pairplot(Negative[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']])
```



Correlation Analysis and Heat map

In [87]:

```
# correlation matrix
dataset_imputed.corr()
```


Out[87]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
Pregnancies	1.000000	0.129459	0.141282	0.082989	0.056027	0.017683	-0.033523	0.544341	0.221898
Glucose	0.129459	1.000000	0.152590	0.182455	0.407699	0.221071	0.137337	0.263514	0.466581
BloodPressure	0.141282	0.152590	1.000000	0.123444	0.045319	0.281805	0.041265	0.239528	0.065068
SkinThickness	0.082989	0.182455	0.123444	1.000000	0.158139	0.480496	0.100966	0.127872	0.215299
Insulin	0.056027	0.407699	0.045319	0.158139	1.000000	0.149468	0.098634	0.136734	0.214411
BMI	0.017683	0.221071	0.281805	0.480496	0.149468	1.000000	0.140647	0.036242	0.292695
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.100966	0.098634	0.140647	1.000000	0.033561	0.173844
Age	0.544341	0.263514	0.239528	0.127872	0.136734	0.036242	0.033561	1.000000	0.238356
Outcome	0.221898	0.466581	0.065068	0.215299	0.214411	0.292695	0.173844	0.238356	1.000000

```
In [88]:
plt.subplots(figsize=(8,8))
sns.heatmap(dataset_imputed.corr(),annot=True)
```

Out[88]:

<AxesSubplot:>



Correlation Results :

- There are not much multicollinearity
- Pregnancies and Age have some positive correlation
- Glucose has some postive corelation with the outcome variable
- Skin thickness and BMI has some positive correlation
- Insulin and Glucose has some positive correlation

Project Task: Week 3 and Week 4 -- Data Modelling and Model Performance Evaluation

Model 1 : Logistic Regression

```
In [89]:
dataset_imputed.head(5)
```

Out[89]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35.00000	155.548223	33.6	0.627	50	1
1	1	85	66	29.00000	155.548223	26.6	0.351	31	0
2	8	183	64	29.15342	155.548223	23.3	0.672	32	1
3	1	89	66	23.00000	94.000000	28.1	0.167	21	0
4	0	137	40	35.00000	168.000000	43.1	2.288	33	1

In [94]:

```
features = dataset_imputed.iloc[:,[0,1,2,3,4,5,6,7]].values
label = dataset_imputed.iloc[:,8].values
```

In [95]:

```
#Train test split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(features,label,test_size=0.2,random_state =10)
```

In [96]:

```
#Create model
from sklearn.linear_model import LogisticRegression
logRegModel = LogisticRegression()
logRegModel.fit(X_train,y_train)
```

Out[96]:

```
LogisticRegression()
```

In [97]:

```
print(logRegModel.score(X_train,y_train))
print(logRegModel.score(X_test,y_test))
```

```
0.7719869706840391
0.7597402597402597
```

In [98]:

```
y_pred = logRegModel.predict(X_test)
print('Accuracy of logistic regression classifier on test set: {:.2f}'.format(logRegModel.score(X_test, y_test)))
```

Accuracy of logistic regression classifier on test set: 0.76

In [99]:

```
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[86  9]
 [28 31]]
```

In [100]:

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.75	0.91	0.82	95
1	0.78	0.53	0.63	59
accuracy			0.76	154
macro avg	0.76	0.72	0.72	154
weighted avg	0.76	0.76	0.75	154


In [101]:

```

from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, logRegModel.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logRegModel.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
print('AUC: %0.3f' % logit_roc_auc)
plt.show()

```

AUC: 0.715



Model 2 : Decision Tree Classifier

In [103]:

```

#Hyper Parameter tuning of max_dept
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics
for i in range(3,20):print("For max_depth = ",i)
DTModel = DecisionTreeClassifier(max_depth=i)
DTModel.fit(X_train,y_train)
y_pred = DTModel.predict(X_test)
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

```

```

For max_depth = 3
For max_depth = 4
For max_depth = 5
For max_depth = 6
For max_depth = 7
For max_depth = 8
For max_depth = 9
For max_depth = 10
For max_depth = 11
For max_depth = 12
For max_depth = 13
For max_depth = 14
For max_depth = 15
For max_depth = 16
For max_depth = 17
For max_depth = 18
For max_depth = 19
Accuracy: 0.7077922077922078

```

Highest Accuracy of Decision Tree Model can be obtained on Max_Depth = 10

In [104]:

```

DTModel = DecisionTreeClassifier(max_depth=10)
DTModel.fit(X_train,y_train)
y_pred = DTModel.predict(X_test)

```

In [105]:

```
DTModel.score(X_train,y_train)
```

Out[105]:

0.9267100977198697

In [106]:

```
DTModel.score(X_test,y_test)
```

Out[106]:

0.7402597402597403

In [107]:

```
print('Accuracy of Decision Tree regression classifier on test set: {:.2f}'.format(DTModel.score(X_test, y_test)))
```

Accuracy of Decision Tree regression classifier on test set: 0.74

In [108]:

```
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[76 19]
 [21 38]]
```

In [109]:

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.80	0.79	95
1	0.67	0.64	0.66	59
accuracy			0.74	154
macro avg	0.73	0.72	0.72	154
weighted avg	0.74	0.74	0.74	154

In [110]:

```
from sklearn.metrics import precision_score
print("Precision score: {}".format(precision_score(y_test,y_pred)))
```

Precision score: 0.6666666666666666

In [111]:

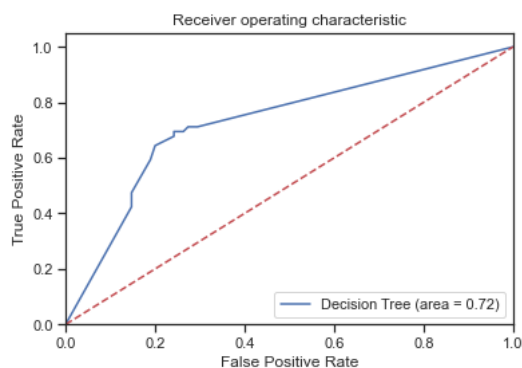
```
from sklearn.metrics import recall_score
print("Recall score: {}".format(recall_score(y_test,y_pred)))
```

Recall score: 0.6440677966101694

In [112]:

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
dt_roc_auc = roc_auc_score(y_test, DTModel.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, DTModel.predict_proba(X_test)[:,:1])
plt.figure()
plt.plot(fpr, tpr, label='Decision Tree (area = {:.2f}) % dt_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('DT_ROC')
print('AUC: {:.3f} % dt_roc_auc)
plt.show()
```

AUC: 0.722



Model 3 : Random Forest Classifier

In [113]:

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
```

In [114]:

```
from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
roc_auc
```

Out[114]:


0.7048171275646743

In [115]:

```
#Hyper Parameter tuning of n_estimators
n_estimators = [1, 2, 4, 8, 16, 32, 64, 100, 200]
train_results = []
test_results = []
for estimator in n_estimators:
    rf = RandomForestClassifier(n_estimators=estimator, n_jobs=-1)
    rf.fit(X_train, y_train)
    train_pred = rf.predict(X_train)
    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
    roc_auc = auc(false_positive_rate, true_positive_rate)
    train_results.append(roc_auc)
    y_pred = rf.predict(X_test)
    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
    roc_auc = auc(false_positive_rate, true_positive_rate)
    test_results.append(roc_auc)
```

In [116]:

```
from matplotlib.legend_handler import HandlerLine2D
line1, = plt.plot(n_estimators, train_results, 'b', label="Train AUC")
line2, = plt.plot(n_estimators, test_results, 'r', label="Test AUC")
plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
plt.ylabel('AUC score')
plt.xlabel('n_estimators')
plt.show()
```



In [117]:

```
rfModel = RandomForestClassifier(n_estimators=60)
rfModel.fit(X_train, y_train)
y_pred = rfModel.predict(X_test)
```

In [118]:

```
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
roc_auc
```

Out[118]:

0.6995539696699375

In [119]:

```
rfModel.score(X_train, y_train)
```

Out[119]:

1.0

In [120]:

```
rfModel.score(X_test,y_test)
```

Out[120]:

0.7402597402597403

In [121]:

```
print('Accuracy of Random Forest regression classifier on test set: {:.2f}'.format(rfModel.score(X_test, y_test)))
```

Accuracy of Random Forest regression classifier on test set: 0.74

In [122]:

```
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[83 12]
 [28 31]]
```

In [123]:

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.75	0.87	0.81	95
1	0.72	0.53	0.61	59
accuracy			0.74	154
macro avg	0.73	0.70	0.71	154
weighted avg	0.74	0.74	0.73	154

In [124]:

```
from sklearn.metrics import precision_score
print("Precision score: {}".format(precision_score(y_test,y_pred)))
```

Precision score: 0.7209302325581395

In [125]:


```
from sklearn.metrics import recall_score
print("Recall score: {}".format(recall_score(y_test,y_pred)))
```

Recall score: 0.5254237288135594

In [126]:

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
rf_roc_auc = roc_auc_score(y_test, rfModel.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, rfModel.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr, tpr, label='Random Forest (area = {:.2f})'.format(rf_roc_auc))
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('RF_ROC')
print('AUC: {:.3f}'.format(rf_roc_auc))
plt.show()
```

AUC: 0.700



Model 4 : Support Vector Machine

In [127]:

```
#Support Vector Classifier
from sklearn.svm import SVC
SVMmodel = SVC(kernel='rbf',
               gamma='auto')
SVMmodel.fit(X_train,y_train)
```

Out[127]:

```
SVC(gamma='auto')
```

In [128]:

```
SVMmodel.score(X_train,y_train)
```

Out[128]:

```
1.0
```

In [129]:

```
SVMmodel.score(X_test,y_test)
```

Out[129]:

```
0.6168831168831169
```

Model 5 : KNN Classifier

In [130]:

```
#Applying K-NN
from sklearn.neighbors import KNeighborsClassifier
knnClassifier = KNeighborsClassifier(n_neighbors=7,
                                   metric='minkowski',
                                   p = 2)
knnClassifier.fit(X_train,y_train)
```

Out[130]:

```
KNeighborsClassifier(n_neighbors=7)
```

In [131]:

```
knnClassifier.score(X_train,y_train)
```

Out[131]:

```
0.8045602605863192
```

In [132]:

```
knnClassifier.score(X_test,y_test)
```

Out[132]:

```
0.7272727272727273
```

In [133]:

```

#Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
# predict probabilities
probs = knnClassifier.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(y_test, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr, fpr, thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")

```

AUC: 0.739


```

True Positive Rate - [0.          0.10169492 0.20338983 0.40677966 0.54237288 0.61016949
 0.77966102 0.93220339 1.          ], False Positive Rate - [0.          0.01052632 0.03157895 0.10526316 0.15789474 0.2736842
1          0.50526316 0.66315789 1.          ] Thresholds - [2.          1.          0.85714286 0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.          ]

```

Out[133]:

Text(0, 0.5, 'True Positive Rate')



In [134]:

```

print('Accuracy of KNN classifier on test set: {:.2f}'.format(knnClassifier.score(X_test, y_test)))

```

Accuracy of KNN classifier on test set: 0.73


In [135]:

```
#Precision Recall Curve for Logistic Regression
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = logRegModel.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = logRegModel.predict(X_test)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y_test, probs)
# calculate F1 score
f1 = f1_score(y_test, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y_test, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.',)
```

f1=0.626 auc=0.759 ap=0.762

Out[135]:

[<matplotlib.lines.Line2D at 0x258c7d75850>]




In [136]:

```
#Precision Recall Curve for KNN
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = knnClassifier.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = knnClassifier.predict(X_test)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y_test, probs)
# calculate F1 score
f1 = f1_score(y_test, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y_test, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.604 auc=0.661 ap=0.624

Out[136]:

[<matplotlib.lines.Line2D at 0x258ca542190>]




In [137]:

```
#Precision Recall Curve for Decision Tree Classifier
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = DTModel.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = DTModel.predict(X_test)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y_test, probs)
# calculate F1 score
f1 = f1_score(y_test, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y_test, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.',)
```

f1=0.655 auc=0.678 ap=0.571

Out[137]:

[<matplotlib.lines.Line2D at 0x258ca20c9a0>]



In [138]:

```


#Precision Recall Curve for Random Forest
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = rfModel.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = rfModel.predict(X_test)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y_test, probs)
# calculate F1 score
f1 = f1_score(y_test, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y_test, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.608 auc=0.698 ap=0.707

Out[138]:

[<matplotlib.lines.Line2D at 0x258ca121280>]



Therefor we observed that Random Forest is best performing model for this dataset ¶

Accuracy of 77%

Precision = 0.78

Recall = 0.54

AUC = 0.70

In []: